

Data Engineering Day 20

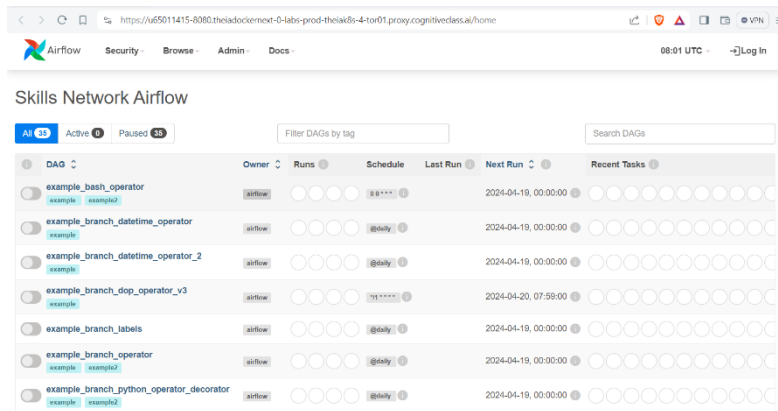
The credit for this course goes to Coursera. [Click More](#)

Another link : [Azure data Engineer](#)

ETL and Data Pipelines with Shell, Airflow and Kafka

Getting started using Apache Airflow UI: [Apache airflow docs](#)

- The figure below represents the UI of Apache airflow.



- to perform etl tasks in airflow, we need to create DAGs and perform shell scriptings for automating the tasks. First of all
- create a file **my_first_dag.sh** copy and paste the code provided below.

```
#!/bin/bash
echo "extract_transform_and_load"

cut -d ":" -f1,3,6 /etc/passwd >
/home/project/airflow/dags/extracted-date.txt
tr ":" "," < /home/project/airflow/dags/extracted-data.txt >
/home/project/airflow/dags/transformed-data.csv
```

- Second step is to create a second file called **my_first_dag.py** and paste the code provided below.

```
- # import the libraries
-
- from datetime import timedelta
- # the DAG object; we will need this to instantiate a DAG
-
- from airflow import DAG
-
- # operators; we need this to write tasks!
-
- from airflow.operators.bash_operator import BashOperator
-
- # this makes scheduling easy
-
- from airflow.utils.dates import days_ago
-
- # defining DAG arguments
-
- # You can override them on a per-task basis during operator
  initializations
-
- default_args = {
-     'owner': 'Ramesh Sannareddy',
-     'start_date': days_ago(0),
-     'email': ['ramesh@somemail.com'],
-     'email_on_failure': False,
-     'email_on_retry': False,
-     'retries': 1,
-     'retry_delay': timedelta(minutes=5),
- }
-
- #DAG arguments are like settings for the DAG.
-
- #The above settings mention
-
- #the owner name,
- #when this DAG should run from: days_ago(0) means today,
- #the email address where the alerts are sent to,
- #whether alert must be sent on failure,
- #whether alert must be sent on retry,
- #the number of retries in case of failure, and
```

```

- #the time delay between retries.
- # defining the DAG
-
- # define the DAG
-
- #A typical DAG definition block looks like below.
- dag = DAG(
-     'my-first-dag',
-     default_args=default_args,
-     description='My first DAG',
-     schedule_interval=timedelta(days=1),
- )
- '''
- Here we are creating a variable named dag by instantiating the
- DAG class with the following parameters.
-
- sample-etl-dag is the ID of the DAG. This is what you see on
- the web console.
-
- We are passing the dictionary default_args, in which all the
- defaults are defined.
-
- description helps us in understanding what this DAG does.
-
- schedule_interval tells us how frequently this DAG runs. In
- this case every day. (days=1).
- '''
-
- # define the task **extract_transform_and_load** to call shell
- script
- #calling the shell script
-
- extract_transform_load = BashOperator(
-     task_id="extract_transform_load",
-     bash_command="/home/project/airflow/dags/my_first_dag.sh ",
-     dag=dag,
- )
-
- '''
- A task is defined using:
-
- A task_id which is a string and helps in identifying the task.

```


What bash command it represents. Here we are calling the shell script `extract_transform_load.sh` which we previously defined
Which dag this task belongs to.

```
'''
# task pipeline
extract_transform_load
```

```
'''


When we execute the task extract_transform_and_load the code in the shell script gets executed.
'''
```

- You will something like below like name of your dags, time, and tasks.

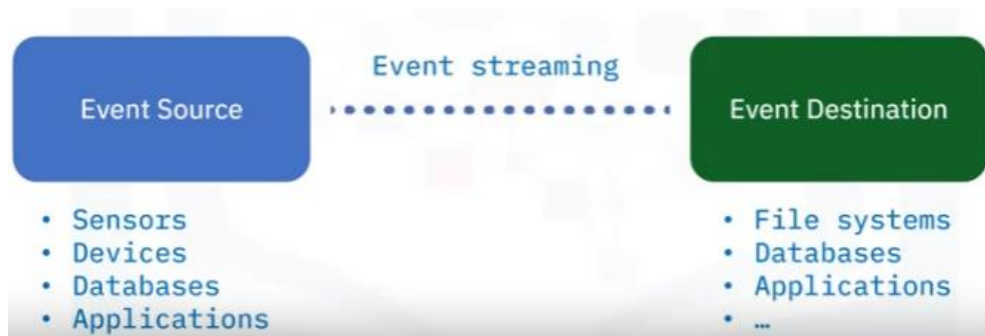


- **Distributed Event Streaming Platform Components:**

- Event is a type of data which describes an entity observable state updates over time.



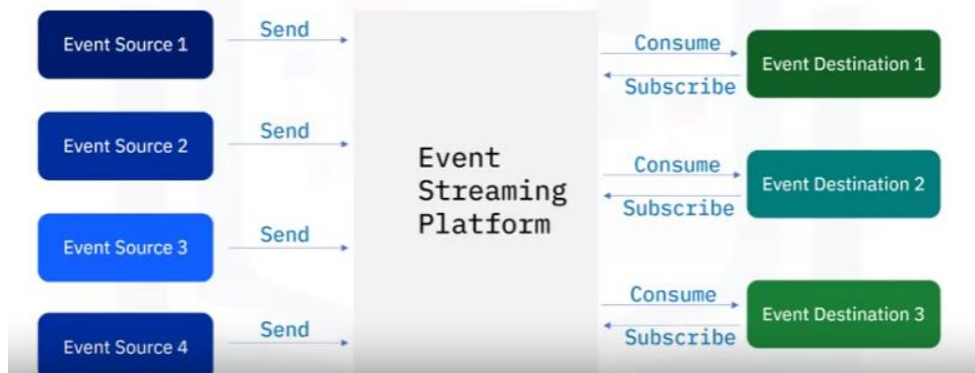
- Common formats of events are primitive, a key value pair, a key value pair with time stamp.
- Event streaming means transformations of events from one destination to another.



Event streaming from multiple event sources to multiple destinations



Event Streaming Platform (ESP)




- **Apache Kafka Overview:**

- Kafka is used to track user activities such as keyboard inputs, mouse, or cursor moments.
- Used for sensor reading, GPS, hardware, and software monitoring. Monitoring
- Used to store logs and financial transitions, data storages, analytics, notifications such as emails.

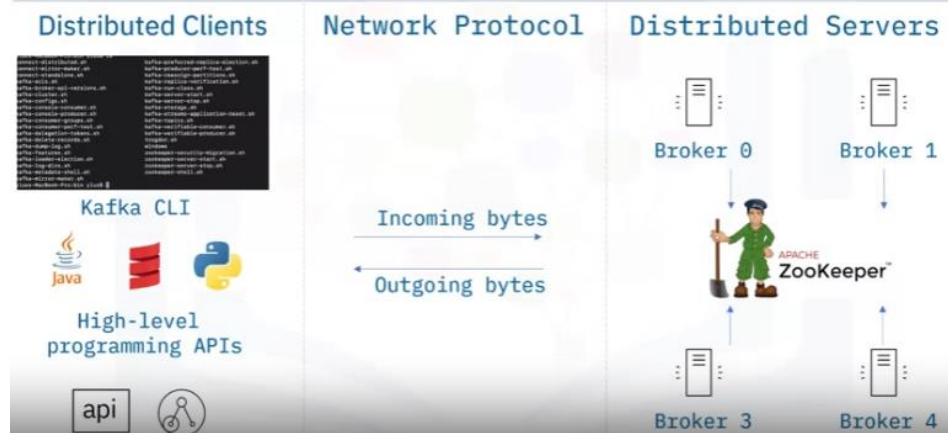
Apache Kafka



Common use cases



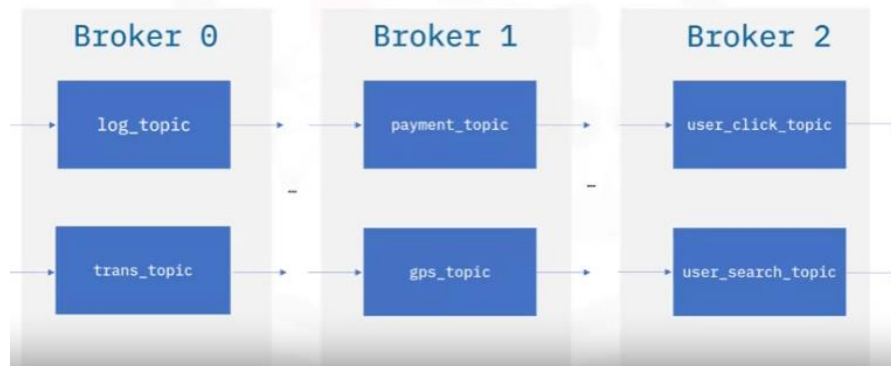
Kafka architecture



- **Building Event Streaming Pipelines using Kafka:**

- Brokers(servers): the dedicated server to receive, store, process and distribute events.

Broker and topic



- Topics: containers or a database of the systems
- Replicants: duplicate partitions into different brokers.
- Partition: divide topics or a database into different brokers or a different server.
- Producers: Kafka client's applications to publish events into topic.

A weather pipeline example

